

OS Lab 7

Sara Mohamed 900203032

Omar Harb 900201063

Mohamed Shaalan 900201539

Roles

Sara Mohamed: *parallel_compute + graphs*

Omar Harb: *parallel_compute + graphs*

Mohamed Shaalan: *sequential_compute + graphs*

Some Dependencies

Pseudocode and description

```
int getcountinfile(char*
filename)
{
    open file;
    int count = 0; int a;
    while (fscanf(a)!=EOF)
        count++;
    close file;
    return count;
}
```

// function to get number of ints in file

```
int* getArray(char* filename, int N)
{
    open file;
    int arr[N];
    int i = 0; int a;
    while (fscanf(a)!=EOF)
    {
        arr[i] = a;
        i++;
    }
    close file;
    return arr;
}
```

// function to save ints from file into array

Sequential_compute

Pseudocode and description

```
double sequential_compute(int
(*f)(int, int), char* filename)
{
    int N =
getcountinfile(filename);
    int* arr = getArray(filename,
N);
    start clock;
    int result = 0; int a, b;

    if (N==0)
        result = 0;
        stop clock;

    else if (N==1)
        result = arr[0];
        stop clock;
```

```
    else
        result = arr[0];
        for (i = 1 to N-1)
            result = f(result, arr[i]);
        stop clock;

    free arr;
    print result to screen;

    return (end clock - start clock);
```

```
}
```

childResult Dependency

Pseudocode and description

*// returns result of computation from
sindex to eindex in arr*

```
int childResult(int* arr, int sindex, int eindex, int (*f)(int, int))
{
    if (sindex == eindex)
        return arr[sindex]; // if start and end index are equal, return the one element

    else if (sindex > eindex)
        return 0; // if end index comes before start index, return 0

    else if (eindex == sindex+1)
        return f(arr[sindex], arr[eindex]); // if two elements, return their computation

    else
        int result = f(arr[sindex], arr[sindex+1]);
        for (int i = sindex+2; i <= eindex; i++)
            result = f(result, arr[i]); // compute result of all elements in interval
        return result;
}
```

Parallel_compute

Pseudocode and description

```
double parallel_compute(int (*f)(int,
int), char* filename, int n_proc)
{
    if (n_proc == 0)
        exit(0); // if n_proc = 0, exit

    int N = getcountinfile(filename);
    int* arr = getArray(filename, N);
    start clock;
    pid_t parentid = getpid();

    int count = 0, index = 0;
    int tempread, tempwrite;
```

```
    int fdesc_arr[n_proc][2];
    int nperprocess = N/n_proc;
    int extraN = nperprocess + (N%n_proc);

    for (i=0 to n_proc-1)
        if (pipe(fdesc_arr[i]) == -1)
            print error message;
            exit(1); // initializing pipes, one per
                    child process
```

```

while (fork() != -1 && count < n_proc)
{
    if (getpid() != parent_id)
        close(fdesc_arr[count][0]);    // closes read side of pipe

    if (count == 0)    // if first element, get result using index and extraN
        tempwrite = childResult(index, index + extraN - 1, f);
        write(fdesc_arr[count][1], &tempwrite, sizeof(tempwrite));
        close(fdesc_arr[count][1]);    // close write side after done
    else    // else, get result using index and nperprocess
        tempwrite = childResult(index, index + nperprocess - 1, f);
        write(fdesc_arr[count][1], &tempwrite, sizeof(tempwrite));
        close(fdesc_arr[count][1]);    // close write side after done
    break;
else
    if (count == 0)
        index += extraN;
    else
        index += nperprocess;
    count++;
}

```

This branch eliminates the need for another pipe- the parent increments the index properly before forking the child that will use that index

```

if (getpid() != parent_id)
    exit(0); // children exit program after their computations

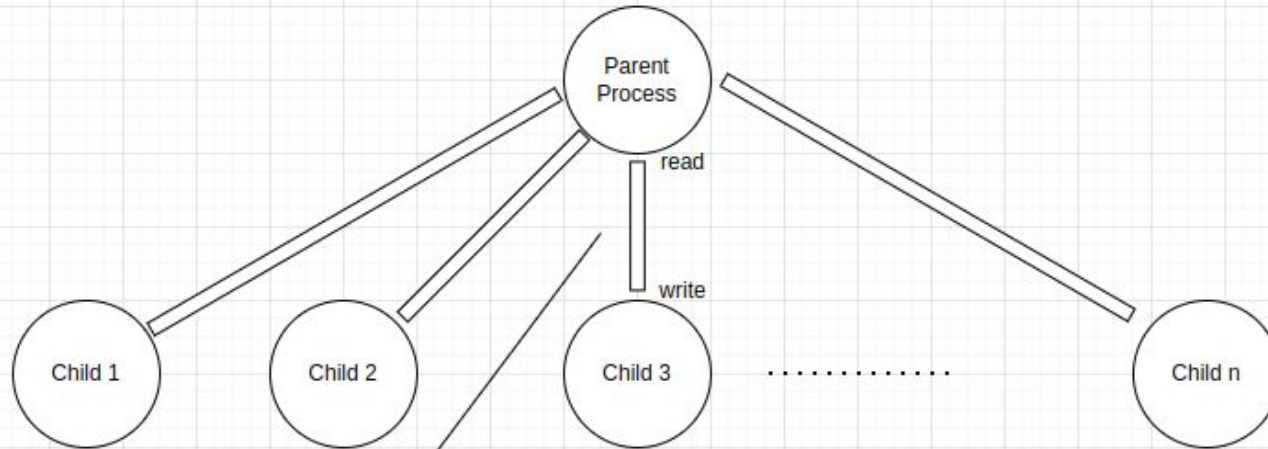
count = 0;
if (getpid() == parent_id)
    close(fdesc_arr[count][1]); // close write side of first pipe
    read(fdesc_arr[count][0], &tempread, sizeof(tempread));
    result = tempread;
    count++; // read first result, store in tempread, increment count

    while (count < n_proc)
        close(fdesc_arr[count][1]);
        read(fdesc_arr[count][1], &tempread, sizeof(tempread));
        result = f(result, tempread);
        Count++; // read remaining results, combine together

end clock;
print result;
free(arr);
return (end clock - start clock);
}

```


Parallel Compute Diagram



A uni-directional pipe where child writes result for parent to read

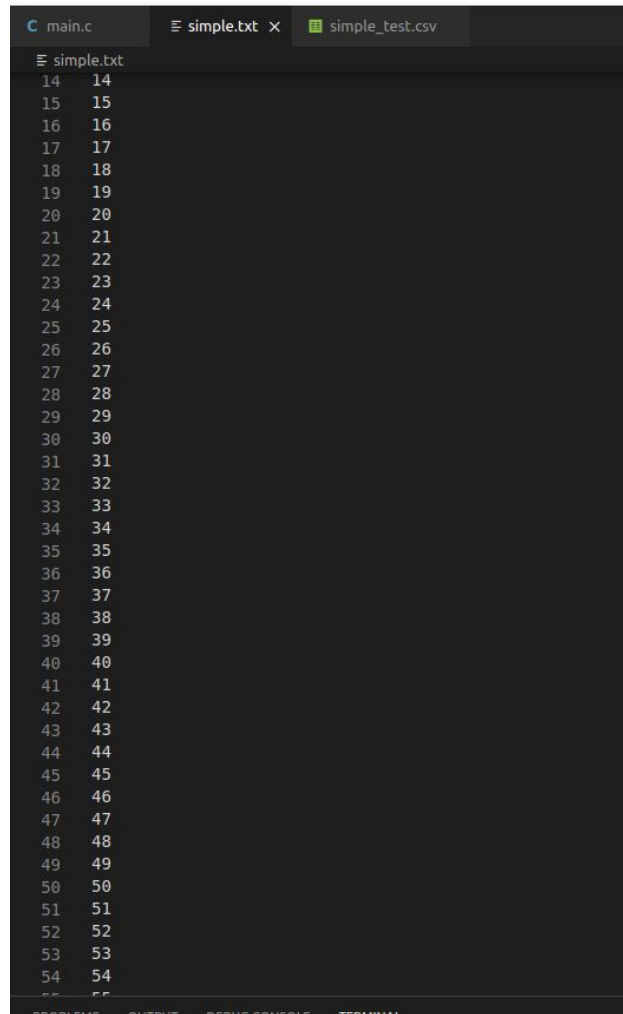
Each child shares a single unidirectional pipe with the parent. The read side is the parent's side, and the write side is the child's. Data (the child's computations) flows from the child to the parent to be combined in result and returned by the function at the end.

Example of Run

```
wrote from 1 to 500 to file 'simple.txt'.  
[sequential compute] 125250, by process 2293023  
    executed in 0.000003 seconds  
[parallel compute with 16 processes]: 125250, by process 2293023  
    executed in 0.000589 seconds
```

This run enters 500 elements (1, 2, 3, ..., 500) to file and then computes them using the two functions

$$\text{result} = \sum_{i=1}^n i = 1 + 2 + 3 + \dots + n = \frac{n(n+1)}{2}$$
$$= 250(501) = 125250.$$



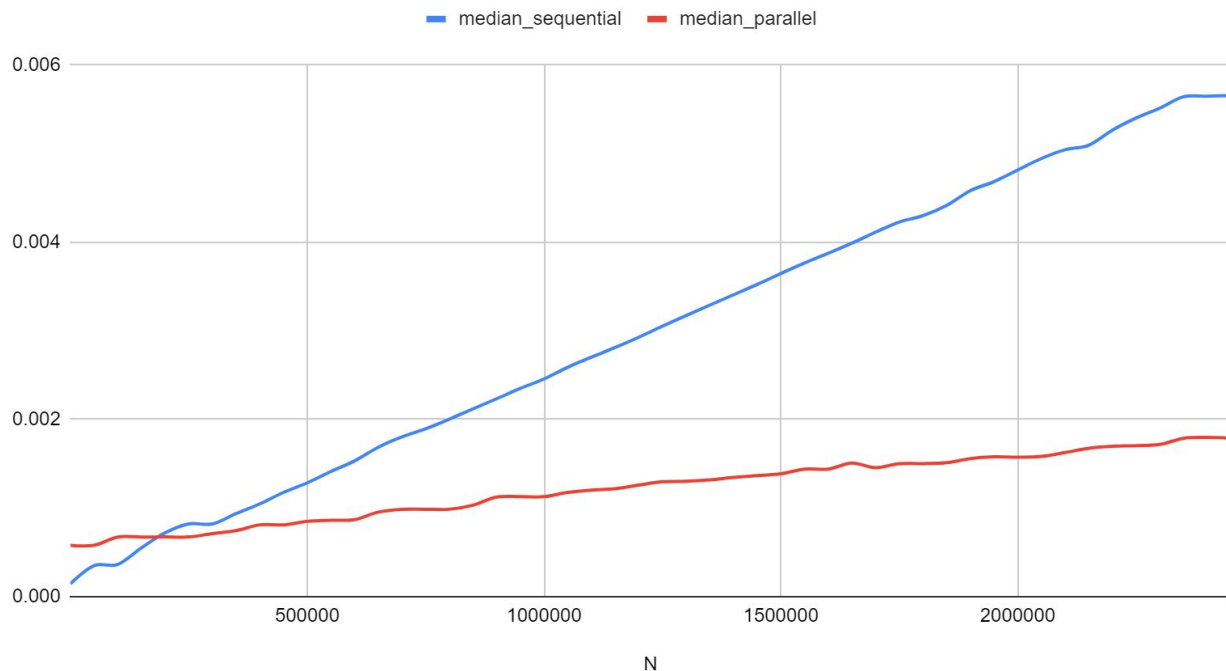
The screenshot shows a code editor with three tabs: 'main.c', 'simple.txt', and 'simple_test.csv'. The 'simple.txt' tab is active, displaying a list of numbers from 14 to 54, with each number appearing twice on the same line (e.g., '14 14', '15 15', etc.). The editor interface includes a sidebar on the left and a bottom panel with tabs for 'PROBLEMS', 'OUTPUT', 'DEBUG CONSOLE', and 'TERMINAL'.

Ranges and Results

This graph contains data ranging from N=1 to N=2.5M, with n_proc set at 16 processes.

	A	B	C	D	E
1	N	sequential	parallel	median_sequential	median_parallel
2	1	0.000001	0.000743	0.00014	0.000575
3	50001	0.00014	0.000575	0.000345	0.000575
4	100001	0.000345	0.000547	0.000359	0.000668
5	150001	0.000359	0.000668	0.000546	0.000668
6	200001	0.000546	0.000806	0.000715	0.000669
7	250001	0.000837	0.000646	0.000816	0.000669
8	300001	0.000715	0.000669	0.000816	0.000707
9	350001	0.000816	0.000707	0.000934	0.000741
10	400001	0.000934	0.000741	0.001043	0.000806
11	450001	0.001043	0.000806	0.001173	0.000806
12	500001	0.001173	0.000845	0.00128	0.000845
13	550001	0.00128	0.000795	0.001408	0.000858
14	600001	0.001408	0.000858	0.001528	0.000864
15	650001	0.001528	0.000864	0.001683	0.00095
16	700001	0.001683	0.00095	0.0018	0.000982
17	750001	0.0018	0.000991	0.001891	0.000982
18	800001	0.001891	0.000982	0.002	0.000982
19	850001	0.002	0.000977	0.002117	0.001028
20	900001	0.002117	0.001028	0.002231	0.00112
21	950001	0.002231	0.001189	0.002348	0.001124
22	1000001	0.002348	0.00112	0.002455	0.001124
23	1050001	0.002455	0.001124	0.002587	0.001172
24	1100001	0.002587	0.001172	0.0027	0.001198
25	1150001	0.0027	0.001198	0.00281	0.001213
26	1200001	0.00281	0.001213	0.002926	0.001255
27	1250001	0.002926	0.001255	0.003051	0.001293
28	1300001	0.003051	0.001299	0.003171	0.001299
29	1350001	0.003171	0.001293	0.003289	0.001315
30	1400001	0.003289	0.001341	0.003406	0.001341
31	1450001	0.003406	0.001315	0.003523	0.001362
32	1500001	0.003523	0.001383	0.003646	0.001383
33	1550001	0.003646	0.001362	0.003766	0.001436
34	1600001	0.003766	0.001436	0.003875	0.001436
35	1650001	0.003875	0.001555	0.003989	0.001504
36	1700001	0.003989	0.001408	0.004114	0.001452
37	1750001	0.004114	0.001504	0.004228	0.001488
38	1800001	0.004228	0.001452	0.0043	0.001488
39	1850001	0.0043	0.001488	0.004415	0.001508
40	1900001	0.004415	0.001508	0.004583	0.001555
41	1950001	0.004583	0.001574	0.004685	0.001574
42	2000001	0.004685	0.001555	0.004816	0.00157
43	2050001	0.004816	0.001579	0.004946	0.001579
44	2100001	0.004946	0.00157	0.005045	0.001623
45	2150001	0.005045	0.001672	0.005096	0.001672
46	2200001	0.005096	0.001623	0.005267	0.001685

graph 1: sequential and parallel (smoothed)

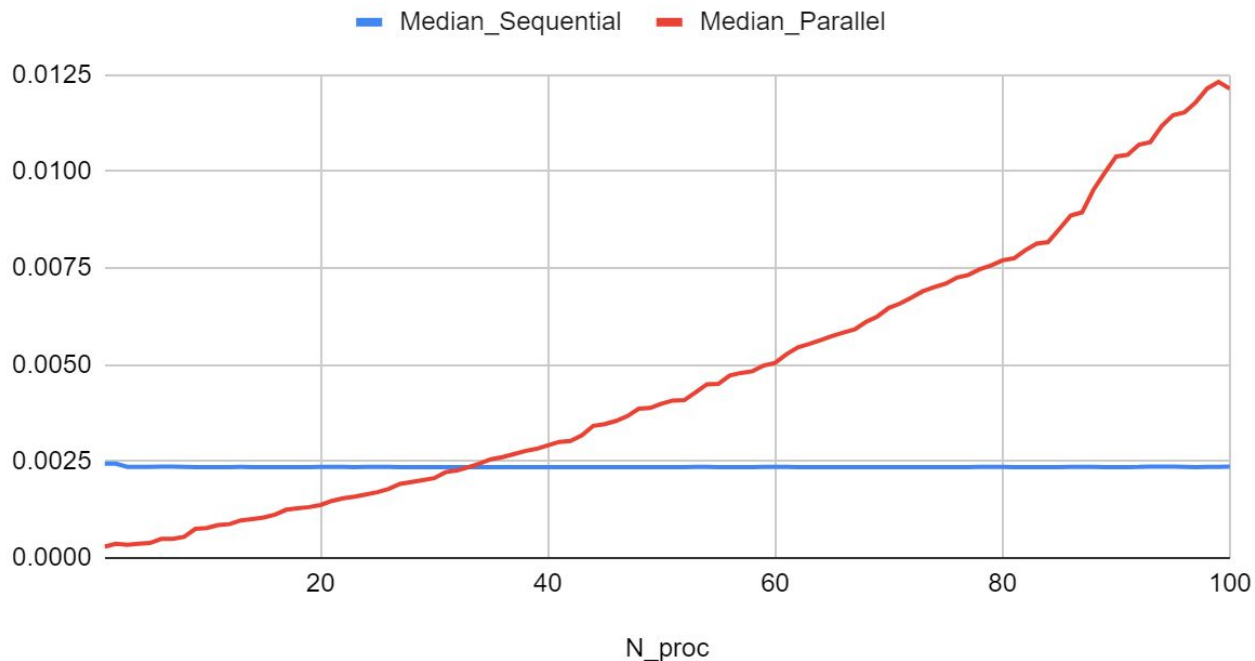


Ranges and Results

This graph contains data ranging from $n_proc=1$ to $n_proc=100$, with N set at 1M data points

	A	B	C	D	E
1	N_proc	Sequential	Parallel	Median_Sequential	Median_Parallel
2	1	0.002349	0.000164	0.00244	0.000292
3	2	0.00244	0.000451	0.00244	0.000366
4	3	0.002442	0.000292	0.002358	0.000339
5	4	0.002358	0.000366	0.002355	0.000366
6	5	0.002351	0.000339	0.002355	0.000388
7	6	0.002355	0.000388	0.002361	0.000491
8	7	0.002361	0.000494	0.002361	0.000494
9	8	0.002361	0.000491	0.002356	0.000545
10	9	0.002356	0.000545	0.002352	0.000753
11	10	0.002348	0.000774	0.002352	0.000774
12	11	0.002352	0.000753	0.002352	0.000849
13	12	0.002352	0.000849	0.002352	0.000876
14	13	0.002355	0.000876	0.002355	0.000972
15	14	0.002349	0.000972	0.002352	0.001008
16	15	0.002357	0.001008	0.002352	0.001049
17	16	0.002352	0.001049	0.002351	0.00112
18	17	0.002348	0.00112	0.002351	0.001249
19	18	0.002351	0.001249	0.002351	0.001285
20	19	0.002351	0.001285	0.002352	0.001317
21	20	0.002352	0.001317	0.002355	0.001337
22	21	0.002361	0.001337	0.002355	0.001475
23	22	0.002355	0.001475	0.002355	0.001544
24	23	0.002352	0.001544	0.002352	0.001589
25	24	0.002355	0.001589	0.002355	0.001644
26	25	0.002351	0.001644	0.002356	0.001702
27	26	0.002356	0.001702	0.002356	0.001788
28	27	0.002356	0.001788	0.00235	0.001917
29	28	0.00235	0.001966	0.002349	0.001966
30	29	0.002349	0.001917	0.002349	0.002016
31	30	0.002349	0.002016	0.002349	0.002065
32	31	0.002351	0.002065	0.002351	0.00222
33	32	0.002348	0.00222	0.002348	0.002254
34	33	0.002351	0.002254	0.00235	0.002346
35	34	0.002346	0.002346	0.00235	0.002442
36	35	0.00235	0.002442	0.00235	0.002553
37	36	0.00235	0.002553	0.00235	0.00261
38	37	0.002351	0.00261	0.002351	0.002685
39	38	0.002347	0.002685	0.002351	0.002766
40	39	0.002362	0.002766	0.002353	0.002824
41	40	0.002351	0.002824	0.002351	0.002913
42	41	0.002353	0.002913	0.002353	0.003003
43	42	0.00235	0.003003	0.00235	0.003029
44	43	0.002353	0.003029	0.002351	0.003175
45	44	0.00235	0.003175	0.002351	0.003419
46	45	0.002351	0.003419	0.002352	0.003463
47	46	0.002351	0.003463	0.002351	0.003513

graph 2: sequential and parallel (smoothed)



Ranges and Results

- Testing and runs were conducted on a PC in the Systems Engineering Lab with 16 CPU cores (hence setting n_proc at 16)
- Smoothing of graphs was conducted by getting the median of elements i.e. element 1 was the *median* of reading 1, 2, and 3; and element 2 was the *median* of reading 2, 3, and 4; etc.
- In the first graph, intersection occurs at approximately $N = 200,000$. Parallel_compute outperforms sequential_compute at that approximate value. In the second graph, intersection occurs at approximately $n_proc = 33 \sim 34$. Sequential compute outperform parallel_compute at that value approximately.