

OS Lab 9

Sara Mohamed 900203032

Omar Harb 900201063

Mohamed Shaalan 900201539

Roles

Omar Harb: threads_compute, bonus functions, results and graphs

Mohamed Shaalan: mmap_compute, bonus functions, results and graphs

Sara Mohamed: threads_compute, bonus functions, results and graphs

Inputs and Outputs of mmap

Inside out `mmap_compute()` function, we initialize the output “result” using the `mmap()` function

```
int* result = mmap(  
    NULL,  
    sizeof(int),  
    PROT_READ | PROT_WRITE,  
    MAP_SHARED | MAP_ANONYMOUS,  
    0,  
    0  
);
```

- `*addr = NULL`, so kernel can place mapping wherever it sees fit
- `length = sizeof(int)`, as result is a shared int
- `prot = PROT_READ | PROT_WRITE`, so all processes sharing this space can read/write
- `flags = MAP_SHARED | MAP_ANONYMOUS`, so all processes mapped to this object share the space; and so that the mapping is anonymous, i.e. not connected to any files
- `fd = 0`
- `offset = 0`, to start from the beginning of the file open on the fd descriptor

Pseudocode of functions

pipes_compute():

```
double pipes_compute(int (*f)(int, int),  
char* filename, int n_proc)  
{
```

```
    if (n_proc == 0)  
        exit;
```

```
    int N = getcountinfile(filename);  
    int* arr = getArray(filename, N);  
    start clock;  
    pid_t parentid = getpid();
```

```
    int count = 0, index = 0;  
    int tempread, tempwrite;
```

```
    int fdesc_arr[n_proc][2];  
    int nperprocess = N/n_proc;  
    int extraN = nperprocess + (N%n_proc);
```

```
    for (i=0 to n_proc-1)  
        if (pipe(fdesc_arr[i]) == -1)  
            print error message;  
            exit(1); // initializing pipes, one per  
                        child process
```

```
pipes_compute():
```

```
while (fork() != -1 && count < n_proc)
{
    if (getpid() != parent_id)
        close(fdesc_arr[count][0]);    // closes read side of pipe

    if (count == 0)    // if first element, get result using index and extraN
        tempwrite = childResult(index, index + extraN - 1, f);
        write(fdesc_arr[count][1], &tempwrite, sizeof(tempwrite));
        close(fdesc_arr[count][1]);    // close write side after done
    else    // else, get result using index and nperprocess
        tempwrite = childResult(index, index + nperprocess - 1, f);
        write(fdesc_arr[count][1], &tempwrite, sizeof(tempwrite));
        close(fdesc_arr[count][1]);    // close write side after done
    break;
else
    if (count == 0)
        index += extraN;
    else
        index += nperprocess;
    count++;
}
```

This branch eliminates the need for another pipe- the parent increments the index properly before forking the child that will use that index

child branch

parent branch

pipes_compute():

```
if (getpid() != parent_id)
    exit; // children exit program after their computations

count = 0;
if (getpid() == parent_id)
    close(fdesc_arr[count][1]); // close write side of first pipe
    read(fdesc_arr[count][0], &tempread, sizeof(tempread));
    result = tempread;
    count++; // read first result, store in tempread, increment count

while (count < n_proc)
    close(fdesc_arr[count][1]);
    read(fdesc_arr[count][1], &tempread, sizeof(tempread));
    result = f(result, tempread);
    count++; // read remaining results, combine together

end clock;
print result;
free arr;
return (end clock - start clock);
}
```

Pseudocode of functions

sequential_compute():

```
double sequential_compute(int
(*f)(int, int), char* filename)
{
    int N =
getcountinfile(filename);
    int* arr = getArray(filename,
N);
    start clock;
    int result = 0;

    if (N==0)
        result = 0;
        stop clock;

    else if (N==1)
        result = arr[0];
        stop clock;

    else
        result = arr[0];
        for (i = 1 to N-1)
            result = f(result, arr[i]);
        stop clock;

    free arr;
    print result to screen;

    return (end clock - start clock);
}
```

Pseudocode of functions

mmap_compute():

```
const int N = getcountinfile(filename);  
int* arr = getArray(filename, N);
```

```
initialize result mmap;
```

```
if(result == MAP_FAILED) {  
    printf("map failed!\n");  
    return 0.0;  
}
```

```
*result = 0;
```

```
Start clock;
```

```
pid_t parent_id = getpid();
```

```
int count = 0, index = 0;
```

```
int nperprocess = N / n_proc;
```

```
int extraN = nperprocess + (N %  
n_proc);
```

*//identical to pipes_compute(), we find the
number of elements each child process will
compute the result of*

mmap_compute():

```
while (fork() != -1 && count < n_proc)
{
    if (getpid() != parent_id)
        if(count == 0) // if first element, get result using index and extraN
            *result = f(*result, childResult(arr, index,
            index+extraN-1, f)); //add to result
        else //else, get result using index and nperprocess
            *result = f(*result, childResult(arr, index,
            index+nperprocess-1, f));
            break;
    else
        if(count == 0)
            index += extraN;
        else
            index += nperprocess;
        count++;
}
```

child branch

parent branch

mmap_compute():

```
if(getpid() != parent_id)
    exit; //children exit
else
    wait for all children to exit;

end clock;
print result;

int err = munmap(result, sizeof(int)); //unmap shared memory
if(err != 0) {
    printf("unmapping failed\n");
    return -1.0;
}

return start clock - end clock;
}
```

Pseudocode of functions

Dependencies: `thread_result()`, `thread_data` struct, global variables

```
void* thread_result(void* arg) {
    thread_data* td = (thread_data*) arg;
    if(td->sindex > td->eindex)
        return NULL;

    else
        pthread_mutex_lock(&mutex1);

        // critical section
        for(int i = td->sindex to td->eindex) {
            tresult = f(tresult, td->arr[i]);

        pthread_mutex_unlock(&mutex1);

}
```

```
// global variables
int tresult = 0;
pthread_mutex_t mutex1;

typedef struct {
    int* arr;
    int sindex;
    int eindex;
} thread_data;
```

Pseudocode of functions

thread_compute():

```
double single_process_thread_compute(int (*f)(int, int), char* filename,  
int n_thread) {
```

```
    int N = getcountinfile(filename);
```

```
    int* arr = getArray(filename, N);
```

```
    start clock;
```

```
    int nperprocess = N / n_thread;
```

```
    int extraN = nperprocess + (N % n_thread);
```

```
    pthread_t tid[n_thread];
```

```
    thread_data tdata[n_thread]; //array of data associated to each thread
```

thread_compute():

```
int error, index = 0;
for(int i = 0 to n_thread - 1) {
    tdata[i].arr = arr; // "assign" each thread the array
    tdata[i].sindex = index; // "assign" start index = index
    if(i == 0)
        tdata[i].eindex = index+extraN-1;
        index += extraN;
    else
        tdata[i].eindex = index+nperprocess-1;
        index += nperprocess;
    error = pthread_create(&(tid[i]), NULL, thread_result, (void*)
&tdata[i]);
    if(error != 0)
        printf("Thread can't be created: [%s]\n", error);
        i--; revert index value; // try again since we have to create n_thread
```

threads

```
}
```

```
for(int i = 0; i < n_thread; i++)
    pthread_join(tid[i], NULL);
// join all threads
```

```
end clock, free arr, print result;
return (end clock - start clock);
```

Run Example

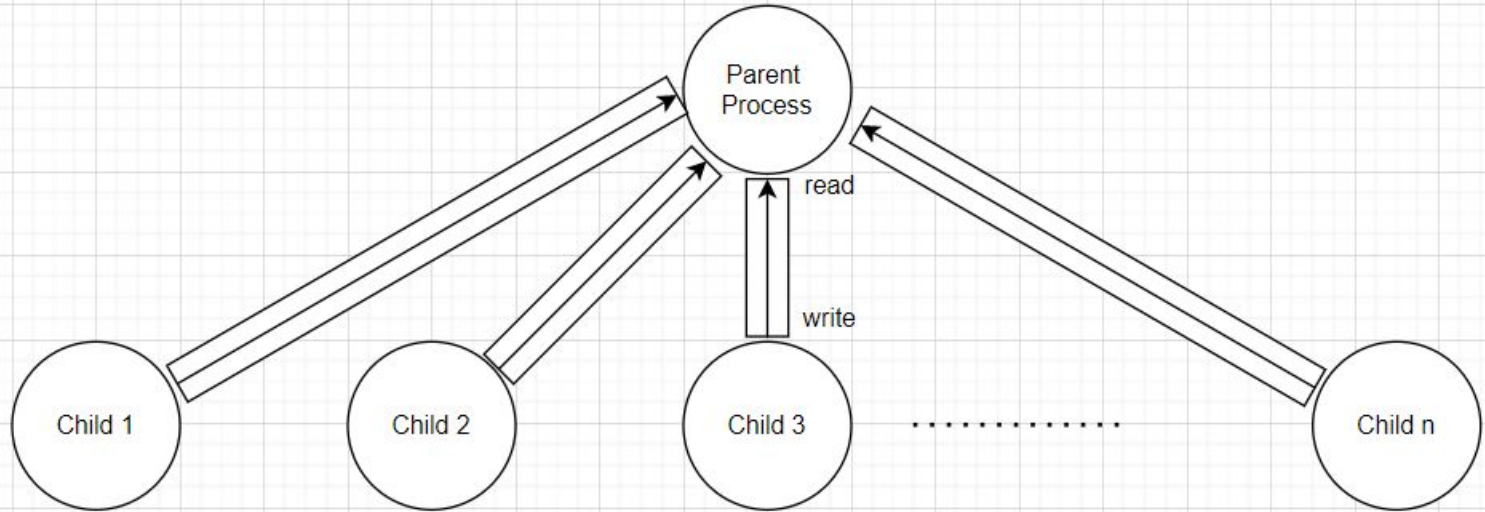
```
Wrote from 1 to 500 to file 'test.txt'.  
[sequential compute] 125250, by process 106826  
    executed in 0.000002 seconds  
[multi pipes compute]: 125250, by process 106826  
    executed in 0.000565 seconds  
[mmap_compute with 16 child processes]: 125250, by process 106826  
    executed in 0.000507 seconds  
[single process thread compute with 16 threads]: 125250, by process 106826  
    executed in 0.000308 seconds
```

This run enters 500 elements (1, 2, 3, ..., 500) to file and then computes them using the four functions


$$\begin{aligned}\text{result} &= \sum_{i=1}^n i = 1 + 2 + 3 + \dots + n = \frac{n(n+1)}{2} \\ &= 250(501) = 125250.\end{aligned}$$

```
14 14  
15 15  
16 16  
17 17  
18 18  
19 19  
20 20  
21 21  
22 22  
23 23  
24 24  
25 25  
26 26  
27 27  
28 28  
29 29  
30 30  
31 31  
32 32  
33 33  
34 34  
35 35  
36 36  
37 37  
38 38  
39 39  
40 40  
41 41  
42 42  
43 43  
44 44  
45 45  
46 46  
47 47  
48 48  
49 49  
50 50  
51 51  
52 52  
53 53  
54 54  
55 55
```

Diagrams: Pipes

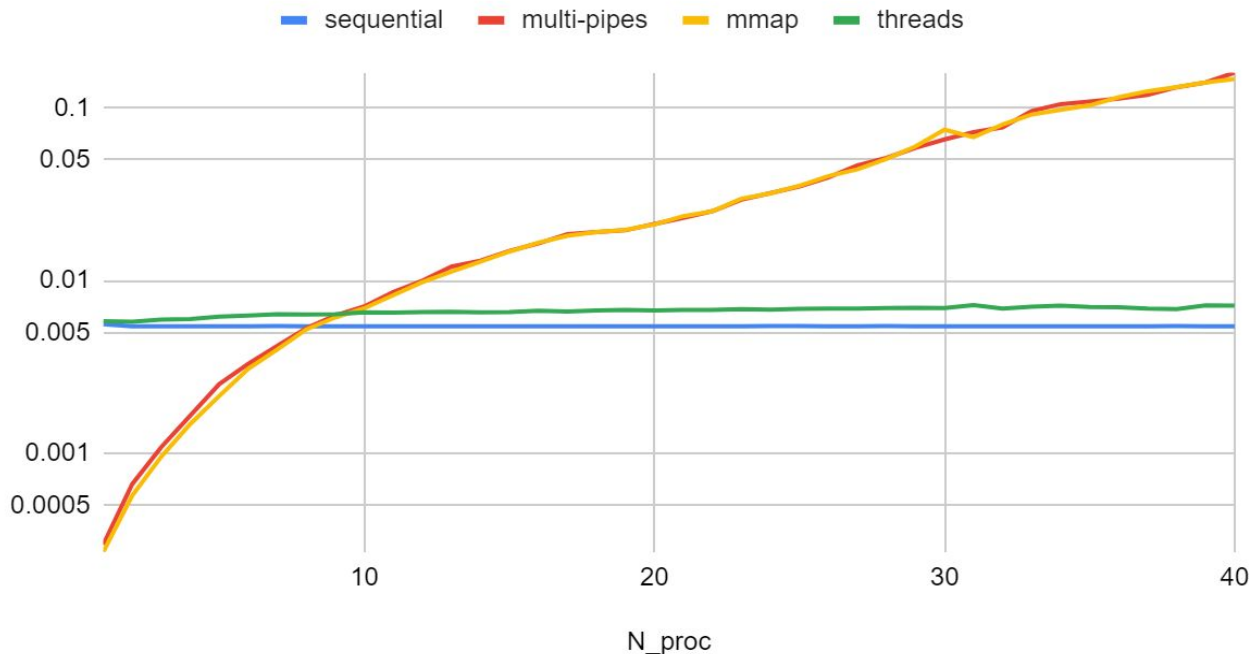


Diagrams: Threads




Graphs

sequential, multi-pipes, mmap and threads



The graph contains data with N fixed at 2M, and n_proc ranging from 1 to 40. Tests were conducted on a PC with 8 CPU cores. Intersection occurs at $n_proc = 8$. Graph presented in log scale for clearer differentiation of results.

Graphs



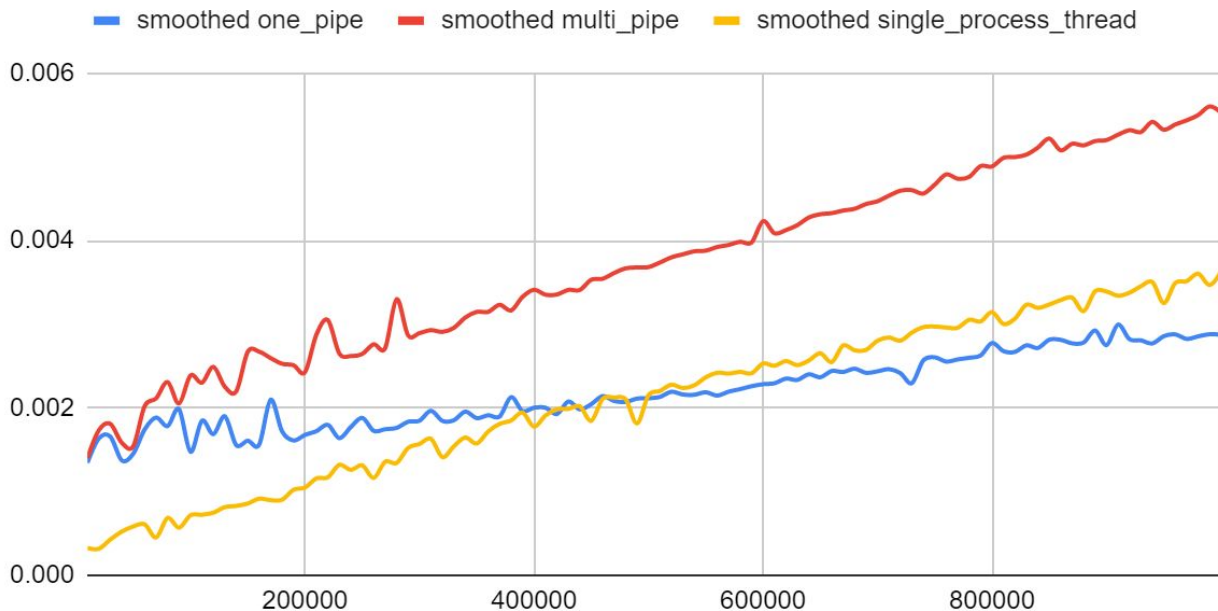
The graph contains data with `n_proc` and `n_thread` fixed at 8, and `N` ranging from 10K to 1M. Tests were conducted on a PC with 8 CPU cores.

Bonus Functions

- `created multi_process_thread_compute, single_process_thread_compute, multi_pipe_compute, and one_pipe_compute` functions (attached in zipped file)
- For comparison graph, fixed `n_thread` and `n_proc` at 8, as testing was conducted on a PC with 8 CPU cores
- `Multi_process_thread_compute` generated error, so did not include it in our generated graph (but still attached in zipped file)

Bonus Graph

N, smoothed one_pipe, smoothed multi_pipe and smoothed single_process_thread



The graph contains data with `n_proc` fixed at 8, and `N` ranging from 10K to 1M. Tests were conducted on a PC with 8 CPU cores.

Ranges and Results

- Testing and runs were conducted on a PC in the Systems Engineering Lab with 8 CPU cores
- Smoothing of graphs was conducted by getting the median of three readings.
- Fixing `n_proc` and `n_thread` to 8, we see that sequential compute always takes the least time, as it has no overhead tied to it, and `thread_compute()` nears it with increasing values of `N`.
- Fixing `N` to `2M`, we see that with `n_proc` and `n_thread` less than 8, sequential and `thread_compute` take longer time, but the behaviour reverses for values on `n_proc` and `n_thread` greater than 8.
- Ranges of tests chosen to both reflect full graph view while maintaining details and information for the smaller values.